

Clarity Health · 3-Minute Demo Runbook

Snowflake Summit · Pediment Booth · for the recorded video walkthrough of the Clarity Health dbt Wizard demo. Looped unattended on the booth screen.

3-5 days MANUAL MODEL BUILD TODAY	90 sec DBT WIZARD BUILD	\$1.2M DENIED CLAIMS AVOIDED / YR	14 hr EARLIER THAN QUALITY COUNCIL
--	--------------------------------------	--	---

ODI architecture · sources, lakehouse, consumers

Fivetran CDC connectors + event streams (Kafka / Kinesis for HL7, telematics, market data)	Iceberg + Glue Open lakehouse on S3, governed catalog, ACID snapshots, time travel	dbt Labs Bronze → Silver → Gold marts, committed as Iceberg snapshots	dbt Wizard Build-time AI · 90s authoring of governed gold marts	Snowflake + 4 Snowflake, Athena, Trino, DuckDB, Spark — same Parquet, no copies
--	--	--	--	--

One set of files on S3. Five engines read the same Parquet through the Glue catalog. Governed end-to-end: HIPAA · row-level access · column masking.

DATA ENGINEERS ANCHOR HERE Architecture · Wizard · Live · Pipeline Show the loop, the four sub-agents, the materialization, the connector deep-links. This is your DE proof.	BUSINESS USERS ANCHOR HERE Executive Cockpit, Patient Registry, Population Health, Quality + Outcomes Show the answers, the KPI tiles, the segments, the executive read. The same governed gold layer powers all of it.
---	--

Frame the demo on business outcomes — never functionality

PAIN Chief Medical Officer can't get answers in time for the Quality Council	REASON Every new question = 3-5 day DE backlog	IMPACT Missed committee, sized exposure, regulator scrutiny	CAPABILITY Build-time AI on the governed gold layer	PROOF Live model build in 90 seconds — on screen
---	---	--	--	---

The mantra: the win is not that dbt Wizard *can* author a model. The win is that the Chief Medical Officer walks into the next Quality Council with a tested, lineage-tracked answer that did not exist this morning. Every beat below ties to that one sentence.

The 3-minute script • five beats

00:00 – 00:25 Beat 1 • The stakeholder hook

WHAT YOU SAY (25S)

"It is Friday morning at Clarity Health. The Chief Medical Officer needs to know: Why did sepsis bundle-compliance dip for the cardiology service line last Tuesday? The Quality Council happens in 14 hours. Today, an answer like this costs three-to-five days of data-engineering work and a missed meeting. We are about to change that."

WHAT YOU SHOW

Open the **Scenario page** (/scenario). The Chief Medical Officer's question is the first thing in frame. Let the T-minus counter tick.

OUTCOME → Anchors the audience on a real business question, not a feature.

00:25 – 00:50 Beat 2 • The cost of the gap

WHAT YOU SAY (25S)

"There is no gold-layer mart that answers this. Three days from now, after a request, a backlog, a scope, a build, and a peer review, the engineering team might land it. The Quality Council will have already happened. Eleven cardiology fallouts means roughly \$1.2 million in denied-claims exposure this quarter alone — and that is just the dollars we can see."

WHAT YOU SHOW

Stay on Scenario. Highlight the four KPI tiles: 3-5 days, 90s, \$1.2M, T-minus to Quality Council.

OUTCOME → Quantifies the cost of the gap. Sets up the 90-second proof.

00:50 – 02:20 Beat 3 • The live build (proof moment • DE-anchor)

WHAT YOU SAY (90S)

"We click Run the Build. Four sub-agents — Explorer, Summary, Worker, Verification — open the dbt project, find the right grain, pull a slice of the gold layer, write the new model, write its tests, and materialize it as an Iceberg snapshot. The viewer sees the SQL appear character by character. By the time we are done talking, the model exists, the tests pass, the lineage graph has updated."

Pause once at: "Worker is anchoring the CTE to the grain the Chief Medical Officer asked for — not a notebook-grade approximation."

And again at: "Verification is writing a uniqueness test on the natural key. This is a governed gold asset, not a SQL snippet."

WHAT YOU SHOW

Click into /wizard-live . Let the playback run at 1x. The agents type SQL and YAML in real time. The step rail moves Discovery → Schema → Inspection → Model Creation → Tests → Materialization. Terminal-window aesthetic.

OUTCOME → Data engineers see the loop end-to-end. The wow is that the output is a tested, governed asset — not a snippet pasted into Slack.

02:20 – 02:45 Beat 4 · The outcome (P&L tie-back · business-anchor)

WHAT YOU SAY (25S)

"The new mart is materialized. The Quality Council convenes in 14 hr, not three days. The Chief Medical Officer walks in with the answer: lactate-redraw orders not auto-firing on shift change for ER-admitted patients between 7-8am — 11 of 14 cardiology fallouts share this signature. That is a workflow fix, not a six-month staffing initiative. And the next time anyone asks this question — Cardiology, Neurology, anywhere — the model is already there."

WHAT YOU SHOW

Cut to /outcome . Materialized model, tests passed, lineage updated, root-cause panel, impact panel. Then quick scroll through the business-user pages (Executive Cockpit, Patient Registry, Population Health, Quality + Outcomes) so the audience sees the same gold layer powering the executive read.

OUTCOME → Closes the loop from question to action AND shows the business-user surface running on the same data.

02:45 – 03:00 Beat 5 · The architectural ask

WHAT YOU SAY (15S)

"Fivetran lands every source. Iceberg on S3 holds the bytes; the Glue catalog governs the snapshots. dbt Labs builds the marts. dbt Wizard authors the gaps. Snowflake reads the gold layer — and so do Athena, Trino, DuckDB, and Spark. One set of files, five engines. That loop closes for clinical-quality questions, denied-claims investigations, readmission deep-dives, and CMS audit prep — all on the same governed lake. Want to see how this maps to your gold layer? Stop by the Pediment booth or scan the QR."

WHAT YOU SHOW

Final card with the architecture loop, QR code to fivetran.com/contact. Tag dbt Labs co-brand.

OUTCOME → Calls out the architectural pattern. Invites a discovery conversation.

Business outcomes to weave through — never lead with features

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
Quality Council gets the answer same-day	3-5 day DE backlog	In time for the meeting, not three days late
\$1.2M exposure sized in 90 seconds	Open gap window	The cost of the gap closed before lunch
Reusable governed mart for every future ask	One-off notebook scripts	The next person who asks gets it for free
CDO does not add headcount to keep up	Hiring DE backfill	Your team stops being the bottleneck
Auditor sees lineage + tests, not a query history	Manual evidence packs	Examiner-ready by default

Anti-patterns · what NOT to say

Don't say: "It runs dbt show against the warehouse and authors a new model file." **Say:** "Ninety seconds from question to materialized answer."

Don't say: "Cortex Analyst." Cortex is intentionally absent from this demo. The story is build-time AI on the governed lake.

Don't say: "dbt Wizard generates SQL." Say: "dbt Wizard authors a governed asset — schema contract, tests, lineage, ownership."

Don't say: "Look at the typing animation." The animation is the texture, not the value.

Don't open with the architecture diagram. Open with the Chief Medical Officer's question. Architecture is the closer.

Don't skip the business pages. After the wizard finishes, scroll through Executive Cockpit, Patient Registry, Population Health, Quality + Outcomes so the audience sees the persona surface that runs on the same gold layer.

Discovery prompts to anchor before the demo

BEFORE YOU PRESS PLAY

- "What questions does your Chief Medical Officer ask that your team cannot answer same-day?"
- "When a stakeholder asks for a new cut of data, what is the average time from ask to merged dbt model?"
- "How many one-off SQL scripts are sitting in someone's Snowflake history that never made it back into the governed layer?"
- "Who is your Economic Buyer for new data investment? Likely audience: CMIO, CDO, COO, or the Quality Office."

AFTER THE DEMO LANDS

- "Which of those agent steps — Discovery, Schema, Inspection, Model Creation, Tests, Materialization — is where your team spends the most time today?"
- "If we sized a pilot on three of your highest-value recurring Chief Medical Officer questions, what would success look like in 90 days?"
- "Who needs to be in a 30-minute architecture review to see how this lands on your gold layer? Audience: CMIO, CDO, COO, or the Quality Office."

Personas at the booth — split the time

DATA ENGINEER · 60% OF BOOTH TIME

Scenario → Live → Outcome → Architecture → Pipeline

Lead with the question, show the four sub-agents, then walk the architecture (Fivetran + Iceberg/Glue + dbt Labs + dbt Wizard + Snowflake/4 engines) and the Pipeline page with real Fivetran connector deep-links.

BUSINESS USER · 40% OF BOOTH TIME

Executive Cockpit, Patient Registry, Population Health, Quality + Outcomes

After the wizard finishes, drop into the dashboard pages. Show that the same governed gold layer powers the executive read. The build-time AI is invisible to them — they see the answer.

Required follow-up materials

- One-pager · Clarity Health · at fivetran-jasonchletsos.github.io/Fivetran-Demo-Repository/one-pager/healthcare-epic/
- Live demo URL · fivetran-jasonchletsos.github.io/Healthcare-EPIC-Snowflake-Demo/
- 3-minute walkthrough video (this recording)
- Architecture diagram on the Architecture tab of the demo
- Cross-industry credibility · Pediment Bank

Booth-floor reminder: the demo is built to run unattended on the booth screen. If a prospect walks up mid-playback, do not restart. Pick up the narration where the playback is — that is also the point. The build finishes whether someone is watching or not. That is the entire pitch.

The gold layer gets acted on, not just queried

Activations is the fourth native NewCo stage, right after Transformations. The instant gold.fct_vital flags an ER-admitted cardiology encounter with lactate drawn, no 3-hour redraw ordered, and less than 30 minutes left on the SEP-1 clock, Activations pushes the alert straight into TigerConnect — no reverse-ETL vendor, no second copy of the data, no analyst in the loop.

TRIGGER 150 min lactate drawn, no 3-hr redraw, ER admit	DESTINATION TigerConnect Care_Team_Channel · alert	BUNDLE COMPLIANCE 79% → 96%+ cardiology SEP-1, targeted
--	---	--

WHAT YOU SAY	WHAT YOU SHOW
"Right after Transformations builds gold.fct_vital, Activations is watching one condition: lactate drawn, no 3-hour redraw ordered, and less than 30 minutes left on the SEP-1 clock."	Toggle the vitals feed so one ER-admitted cardiology patient crosses the 150-minute mark live on screen.
"The moment that flag flips true in the governed gold layer, Activations pushes patient_mrn, encounter_id, service_line, and minutes_to_bundle_deadline straight into TigerConnect — no custom API, no nightly batch, no analyst in the loop."	Architecture page → Activations panel → /activations-live: Segment Definition → Field Mapping → API Push.
"The nurse gets the alert before the deadline closes, not a root-cause report after it."	/activations-live → Destination Confirmation table, alert card rendered with the matching patient and deadline fields.
"That's the difference between the eleven fallouts Dr. Nair's team found last week and catching case twelve before it happens. One governed flag, one sync, zero code."	Before/After panel → Reverse-ETL hop row: Census + Hightouch replaced by NewCo Activations, native.

Crisis-week baseline: 11 of 14 cardiology SEP-1 bundle fallouts traced to the 07:00–08:00 shift-change redraw gap. Closing the pre-deadline window is modeled to avoid \$18k–\$34k in excess length-of-stay cost per averted case and keep SEP-1 out of the CMS Hospital VBP penalty band. Synthetic data — not for clinical use.